

EE2211 Pre-Tutorial 5

Dr Feng LIN

feng_lin@nus.edu.sg



Agenda

- Recap
- Self-learning
- Tutorial 5



Recap

- Functions, Derivative and Gradient
 - Inner product, linear/affine functions
 - Maximum and minimum, partial derivatives, gradient
- Least Squares, Linear Regression
 - Objective function, loss function
 - Least square solution, training/learning and testing/prediction
 - Linear regression with multiple outputs

Linear and Affine Functions

Linear Functions

A function $f: \mathcal{R}^d \rightarrow \mathcal{R}$ is **linear** if it satisfies the following **two properties**:

- **Homogeneity** $f(\alpha \mathbf{x}) = \alpha f(\mathbf{x})$ **Scaling**
- **Additivity** $f(\mathbf{x} + \mathbf{y}) = f(\mathbf{x}) + f(\mathbf{y})$ **Adding**

Inner product function

$$f(\mathbf{x}) = \mathbf{a}^T \mathbf{x} = a_1 x_1 + a_2 x_2 + \cdots a_d x_d$$

Linear function is affine, but affine not necessarily to be linear function


Affine function

$$f(\mathbf{x}) = \mathbf{a}^T \mathbf{x} + b \quad \text{scalar } b \text{ is called the offset (or bias)}$$

Functions: Maximum and Minimum

- $f(x)$ has a **local minimum** at $x = c$ if $f(x) \geq f(c)$ for every x in some open interval around $x = c$
- $f(x)$ has a **global minimum** at $x = c$ if $f(x) \geq f(c)$ for all x in the domain of f

A local and a global minima of a function



Note: An **interval** is a set of real numbers with the property that any number that lies between two numbers in the set is also included in the set.

An **open interval** does not include its endpoints and is denoted using parentheses. E.g. $(0, 1)$ means “all numbers greater than 0 and less than 1”.

Functions: Maximum and Minimum

Max and Arg Max

- Given a set of values $\mathcal{A} = \{a_1, a_2, \dots, a_m\}$,
- The operator $\max_{a \in \mathcal{A}} f(a)$ returns the highest value $f(a)$ for all elements in the set \mathcal{A}
- The operator $\arg \max_{a \in \mathcal{A}} f(a)$ returns the element of the set \mathcal{A} that maximizes $f(a)$
- When the set is **implicit** or **infinite**, we can write


$$\max_a f(a) \quad \text{or} \quad \arg \max_a f(a)$$

E.g. $f(a) = 3a, a \in [0,1] \rightarrow \max_a f(a) = 3$ and $\arg \max_a f(a) = 1$

Min and **Arg Min** operate in a similar manner

Note: **arg max** returns a value from the **domain** of the function and **max** returns from the **range (codomain)** of the function.

Derivative and Gradient



- The **derivative** f' of a function f is a function that describes how fast f grows (or decreases)
 - If the derivative is a **constant** value, e.g. 5 or -3
 - The function f grows (or decreases) constantly at any point x of its domain
 - When the derivative f' is a function
 - If f' is **positive** at some x , then the function f **grows** at this point
 - If f' is **negative** at some x , then the function f **decreases** at this point
 - The derivative of **zero** at x means that the function's **slope** at x is **horizontal** (e.g. **maximum** or **minimum** points)
- The process of finding a derivative is called **differentiation**.
- **Gradient** is the generalization of derivative for functions that take several inputs (or one input in the form of a vector or some other complex structure).

Derivative and Gradient

The gradient of a function is a vector of **partial derivatives**

Differentiation of a **scalar** function w.r.t. a **vector**

If $f(\mathbf{x})$ is a **scalar function** of d variables, \mathbf{x} is a $d \times 1$ vector.

Then differentiation of $f(\mathbf{x})$ w.r.t. \mathbf{x} results in a $d \times 1$ vector

$$\frac{df(\mathbf{x})}{d\mathbf{x}} = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_d} \end{bmatrix}$$

This is referred to as the **gradient** of $f(\mathbf{x})$ and often written as $\nabla_{\mathbf{x}} f$.

Derivative and Gradient

Partial Derivatives

Differentiation of a **vector** function w.r.t. a **vector**

If $\mathbf{f}(\mathbf{x})$ is a **vector function** of size $h \times 1$ and \mathbf{x} is a $d \times 1$ vector.
Then differentiation of $\mathbf{f}(\mathbf{x})$ results in a $h \times d$ matrix

$$\frac{d\mathbf{f}(\mathbf{x})}{d\mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \dots & \frac{\partial f_1}{\partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_h}{\partial x_1} & \dots & \frac{\partial f_h}{\partial x_d} \end{bmatrix}$$

The matrix is referred to as the **Jacobian** of $\mathbf{f}(\mathbf{x})$

Derivative and Gradient

Some Vector-Matrix Differentiation Formulae

$$\frac{d\mathbf{Ax}}{d\mathbf{x}} = \mathbf{A}$$

$$\frac{d(\mathbf{b}^T \mathbf{x})}{d\mathbf{x}} = \mathbf{b} \qquad \frac{d(\mathbf{y}^T \mathbf{Ax})}{d\mathbf{x}} = \mathbf{A}^T \mathbf{y}$$

$$\frac{d(\mathbf{x}^T \mathbf{Ax})}{d\mathbf{x}} = (\mathbf{A} + \mathbf{A}^T) \mathbf{x}$$

$$f(\mathbf{x}) = \mathbf{a}^T \mathbf{x} = a_1 x_1 + a_2 x_2 + \cdots a_d x_d$$

Derivations: <https://www.math.uwaterloo.ca/~hwolkowi/matrixcookbook.pdf>

Linear Regression

- **Linear regression** is a popular regression learning algorithm that learns a model which is a linear combination of features of the input example.

$$\mathbf{X}\mathbf{w} = \mathbf{y}, \quad \mathbf{X} \in \mathcal{R}^{m \times d}, \quad \mathbf{w} \in \mathcal{R}^{d \times 1}, \quad \mathbf{y} \in \mathcal{R}^{m \times 1}$$

$$\mathbf{X} = \begin{bmatrix} x_{1,1} & x_{1,2} & \cdots & x_{1,d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{m,1} & x_{m,2} & \cdots & x_{m,d} \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_d \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

Linear Regression

Problem Statement: To predict the unknown y for a given x (**testing**)

- We have a collection of labeled examples (**training**) $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$
 - m is the size of the collection
 - \mathbf{x}_i is the d -dimensional feature vector of example $i = 1, \dots, m$ (input)
 - y_i is a real-valued target (1-D)
 - Note:
 - when y_i is **continuous** valued, it is a **regression problem**
 - when y_i is **discrete** valued, it is a **classification problem**
- We want to **build a model** $f_{\mathbf{w},b}(\mathbf{x})$ as a linear combination of features of example \mathbf{x} : $f_{\mathbf{w},b}(\mathbf{x}) = \mathbf{x}^T \mathbf{w} + b$
where \mathbf{w} is a d -dimensional vector of parameters and b is a real number.
- The notation $f_{\mathbf{w},b}$ means that the model f is parametrized by two values: \mathbf{w} and b


Linear Regression

Learning objective function

- To find the optimal values for \mathbf{w}^* and b^* which **minimizes** the following expression:

$$\frac{1}{m} \sum_{i=1}^m (f_{\mathbf{w},b}(\mathbf{x}_i) - y_i)^2$$

- In mathematics, the expression we minimize or maximize is called an **objective function**, or, simply, an **objective**



$(f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$ is called the **loss function**: a measure of the difference between $f_{\mathbf{w}}(\mathbf{x}_i)$ and y_i or a penalty for misclassification of example i .

Linear Regression

Learning objective function (using simplified notation hereon)

- To find the optimal values for \mathbf{w}^* which **minimizes** the following expression:

$$\sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$$

with $f_{\mathbf{w}}(\mathbf{x}_i) = \mathbf{x}_i^T \mathbf{w}$,

where we define $\mathbf{w} = [b, w_1, \dots, w_d]^T = [w_0, w_1, \dots, w_d]^T$,

and $\mathbf{x}_i = [1, x_{i,1}, \dots, x_{i,d}]^T = [x_{i,0}, x_{i,1}, \dots, x_{i,d}]^T, i = 1, \dots, m$

- This particular choice of the loss function is called **squared error loss**

Note: The normalization factor $\frac{1}{m}$ can be omitted as it does not affect the optimization.

Linear Regression

$$\sum_{i=1}^m (f_{\mathbf{w},b}(\mathbf{x}_i) - y_i)^2$$

- All model-based learning algorithms have a **loss function**
- What we do to find the best model is **to minimize the objective** known as the **cost function**
- **Cost function** is a sum of **loss functions** over training set plus possibly some model complexity penalty (regularization)
- In linear regression, the cost function is given by the *average loss*, also called the **empirical risk** because we do not have all the data (e.g. testing data)
 - The average of all penalties is obtained by applying the model to the training data

Linear Regression

Learning (Training)

- Consider the set of feature vector \mathbf{x}_i and target output y_i indexed by $i = 1, \dots, m$, a linear model $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^T \mathbf{w}$ can be stacked as

$$\begin{aligned} f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w} & \longleftrightarrow \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix} \\ & \text{Learning Model} \qquad \qquad \qquad \text{Learning target vector} \\ & = \begin{bmatrix} \mathbf{x}_1^T \mathbf{w} \\ \vdots \\ \mathbf{x}_m^T \mathbf{w} \end{bmatrix} \\ \text{where } \mathbf{x}_i^T \mathbf{w} &= [1, x_{i,1}, \dots, x_{i,d}] \begin{bmatrix} b \\ w_1 \\ \vdots \\ w_d \end{bmatrix} \end{aligned}$$

Note: The **bias/offset term** is responsible for **translating** the line/plane/hyperplane away from the origin.

Linear Regression

Least Squares Regression

In vector-matrix notation, the minimization of the objective function can be written compactly using $\mathbf{e} = \mathbf{X}\mathbf{w} - \mathbf{y}$:

$$\begin{aligned} J(\mathbf{w}) &= \mathbf{e}^T \mathbf{e} \\ &= (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= (\mathbf{w}^T \mathbf{X}^T - \mathbf{y}^T) (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y} \\ &= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}. \end{aligned}$$

Note: when $f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w}$, then

$$\sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}).$$

Linear Regression

Differentiating $J(\mathbf{w})$ with respect to \mathbf{w} and setting the result to $\mathbf{0}$:

$$\begin{aligned}\frac{\partial}{\partial \mathbf{w}} J(\mathbf{w}) &= \mathbf{0} \\ \frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}) &= \mathbf{0} \\ \Rightarrow 2\mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{X}^T \mathbf{y} &= \mathbf{0} \\ \Rightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} &= \mathbf{X}^T \mathbf{y}\end{aligned}$$

\Rightarrow Any minimizer $\hat{\mathbf{w}}$ of $J(\mathbf{w})$ must satisfy $\mathbf{X}^T (\mathbf{X} \mathbf{w} - \mathbf{y}) = \mathbf{0}$.

If $\mathbf{X}^T \mathbf{X}$ is invertible, then


Learning/training:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Prediction/testing:

$$\hat{f}_{\mathbf{w}}(\mathbf{X}_{new}) = \mathbf{X}_{new} \hat{\mathbf{w}}$$

Linear Regression



$$\mathbf{X}\mathbf{w} = \mathbf{y}$$

$$\frac{1}{m} \sum_{i=1}^m (f_{\mathbf{w},b}(\mathbf{x}_i) - y_i)^2$$

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$\hat{f}_{\mathbf{w}}(\mathbf{X}_{new}) = \mathbf{X}_{new} \hat{\mathbf{w}}$$

- \mathbf{X} : Samples
- \mathbf{y} : Target values

- Linear or Affine function
- Squared error loss function

- Check the invertibility
- Least square approximation (left-inverse)

- Prediction for new inputs
- Testing: Mean Squared Error (MSE)

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Linear Regression

Learning of Vectored Function (Multiple Outputs)

For one sample: a linear model $\mathbf{f}_w(\mathbf{x}) = \mathbf{x}^T \mathbf{W}$

Vector function

For m samples: $\mathbf{F}_w(\mathbf{X}) = \mathbf{X}\mathbf{W} = \mathbf{Y}$

$$\begin{array}{l} \text{Sample 1} \longrightarrow \\ \vdots \\ \text{Sample } m \longrightarrow \end{array} \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_m^T \end{bmatrix} \mathbf{W} = \begin{bmatrix} 1 & x_{1,1} & \dots & x_{1,d} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{m,1} & \dots & x_{m,d} \end{bmatrix} \underbrace{\begin{bmatrix} w_{0,1} & \dots & w_{0,h} \\ w_{1,1} & \dots & w_{1,h} \\ \vdots & \ddots & \vdots \\ w_{d,1} & \dots & w_{d,h} \end{bmatrix}}_h$$

$$\begin{array}{l} \text{Sample 1's output} \longrightarrow \\ \vdots \\ \text{Sample } m\text{'s output} \longrightarrow \end{array} \begin{bmatrix} y_{1,1} & \dots & y_{1,h} \\ \vdots & & \vdots \\ y_{m,1} & \dots & y_{m,h} \end{bmatrix} \underbrace{\quad}_h \quad m$$

$$\mathbf{X} \in \mathcal{R}^{m \times (d+1)}, \mathbf{W} \in \mathcal{R}^{(d+1) \times h}, \mathbf{Y} \in \mathcal{R}^{m \times h}$$

Linear Regression

Objective: $\sum_{i=1}^m (\mathbf{f}_{\mathbf{w}}(\mathbf{x}_i) - \mathbf{y}_i)^2 = \mathbf{E}^T \mathbf{E}$

Least Squares Regression of Multiple Outputs

In matrix notation, the sum of squared errors cost function can be written compactly using $\mathbf{E} = \mathbf{XW} - \mathbf{Y}$:

$$\begin{aligned} J(\mathbf{W}) &= \text{trace}(\mathbf{E}^T \mathbf{E}) \\ &= \text{trace}[(\mathbf{XW} - \mathbf{Y})^T (\mathbf{XW} - \mathbf{Y})] \end{aligned}$$

If $\mathbf{X}^T \mathbf{X}$ is invertible, then

Learning/training: $\hat{\mathbf{W}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$ \mathbf{Y} is a matrix

Prediction/testing: $\hat{\mathbf{f}}_{\mathbf{w}}(\mathbf{X}_{new}) = \mathbf{X}_{new} \hat{\mathbf{W}}$



THANK YOU